

Liberta: A Hybrid Transformer-Linguistic Model for Machine-generated Text Detection

Fabian Kobylanskyj, Barış Şimşek, Jian Sun and Jiwoo Lee
Technical University of Munich, Germany

Abstract

With the rapid advancement of Large Language Models (LLMs) capable of generating text highly resembling human writing, a critical task is to distinguish between human written and machine-generated text. This paper describes our work as part of the NLP Bachelor Praktikum SS2025 on the SemEval 2024 Task 8: *Multidomain, Multimodel and Multilingual Machine-Generated Text Detection*. Our method focused on combining RoBERTa embeddings with linguistic and token probabilistic features on tackling Subtask A: Human vs Machine Text Classification. The final model’s architecture was systematically developed on the monolingual track to obtain a reliably working variant, which later was adapted and evaluated in a multilingual context. Our results demonstrate that this hybrid approach can compete with the top performing teams of the competition, achieving an accuracy of 88.64% on the official test dataset for the monolingual, and an accuracy of x% for the multilingual track.

1 Introduction

LLMs have a significant influence on our daily tasks across a wide range of domains, from education and research to content creation and communication. Since their impact on society is significantly shaping how we live and work, the potential risks associated with them are equally substantial. From spreading misinformation (Kreps et al., 2022), gaining unfair advantages in exams to plagiarism (Wahle et al., 2022), and a societal risk when it becomes unclear if communication or information comes from a human or machine, there are many dangers and potentially malicious use cases that make it crucial to develop automated means for detecting machine-generated text. In response, different detection tools have been developed. However, since a great deal of effort is put into developing and improving LLMs, the detection of machine-generated text turned out to be an

ongoing arms race where the rapid improvement of LLMs challenges detectors to deliver accurate results. A striking example for this is the OpenAI classifier, which was released in 2023 and discontinued shortly after due to a low rate of accuracy (Jan Hendrik Kirchner, 2023). The SemEval 2024 Task 8: *Multidomain, Multimodel and Multilingual Machine-Generated Text Detection* addresses this challenge and provides an incentive to find solutions to the detection of machine-generated text across multiple generators (Wang et al., 2024). In the following we describe our solution to the monolingual track of Subtask A: Human vs Machine Text Classification, which is a hybrid model that combines RoBERTa embeddings, readability features and token-level probabilistic features. We also present and discuss the adaption of this hybrid model to the multilingual track of Subtask A, which additionally to English includes Arabic, German and Italian as surprising languages. During the development of our models we adhered to the same conditions as participating teams in the shared task in order to ensure a fair comparison. Our results suggest that a hybrid model that combines those three feature sets can significantly improve classification performance and yield a more robust approach than relying on a PLM alone. While RoBERTa successfully captures semantic nuances of the input text, the engineered features provide complementary information that RoBERTa alone misses.

2 Related Work

For the binary classification task of machine-generated text detection as in Subtask A, there are two main approaches. Zero-shot methods and supervised methods (Wang et al., 2024). Zero-shot methods have the advantage of not relying on curated datasets of human written and machine-generated text. This can be particularly beneficial




Figure 1: Overview of the proposed Hybrid Model Architecture combining three different feature vectors.

in the context of the previously mentioned "arms race", where newer LLMs are already far ahead of their predecessors. Although, zero-shot detectors are considered less accurate, innovations like DetectGPT, which only uses log probabilities computed by the model of interest, show that they can accelerate the development of more robust and generalizable methods of machine-generated text detection (Mitchell et al., 2023).

In the shared task, however, most prominently used by the participating teams were supervised training methods. Among those, there are three main groups of methods. The first one is a fine-tuned detection, which uses a pre-trained language model (PLM) and further trains it on the training dataset. Team USTC-BUPT is among the best performing teams at the monolingual track of Subtask A only utilizing PLMs. By adding a gradient reversal layer and a domain classifier that additionally outputs domain labels they achieved an accuracy of 96.10% (Guo et al., 2024).

A second major line focuses on statistical and probabilistic methods. These methods build on the principle that LLMs choose tokens they assign a higher probability to, while humans tend to write in a more "surprising" way. The first ranking Team Genaios employ their model LLMixtic, which mixes token-level probabilistic features from four different Llama-2 models forming a feature vector that is fed into a Transformer Encoder. Their method proved highly effective, achieving an accuracy of 96.88% on the monolingual and 75.67% on the multilingual track (Sarvazyan et al., 2024). Finally, there is a stylometric detection that analyzes linguistic features to distinguish between machine-generated and human written text. By testing a variety of different features as lexical diversity, text statistics and readability features and combining

them with CLS embeddings from a transformer encoder model, Team PetKaz achieved an accuracy of 90.51% on the monolingual track (Petukhova et al., 2024).

3 Methodology

Figure 1 depicts the architecture of our best performing model. It combines three distinct types of features extracted from an input text. The semantic embeddings of a fine-tuned auto-encoder model, different readability metrics and token-level probabilistic features. This feature vector is passed through a feed-forward neural network for classification.

3.1 Semantic Feature Extraction

We employ RoBERTa-base, fine-tuned on the Machine-generated text classification task. The CLS Token represents the entire input sequence and is placed at the beginning of the input, capturing important semantic and contextual information. For each input we extract the CLS Token embedding, which for RoBERTa-base is 768-dimensional, and use it as the first part of our feature vector.

3.2 Readability Feature Extraction

While today’s LLMs generate grammatically correct and coherent text, their output subtly differs from human written text, which we tried to make use of by devoting another part of our feature vector to a detailed linguistic analysis of the input sequence. Team Petkaz at the shared task, whose architecture strongly resembles ours, suggests a range of linguistic features such as text statistics like the number of difficult words, stylometric features, lexical diversity and readability features. Based on our data analysis, we decided to include the latter into our model. Previous works suggest

that machine-generated text tends to be more difficult to read, i.e more complex than human written equivalents. (Georgiou, 2024) for example, reports higher Flesch Reading Ease scores for human texts, which suggest that they are easier to read. Guided by this thesis, we calculate a variety of readability metrics including the following.

Flesch Reading Ease and Flesch-Kincaid Grade Level. Focuses on sentence length and syllable count per word.

$$FK = 0.39 \times \frac{\text{total words}}{\text{total sentences}} + 11.8 \times \frac{\text{total syllables}}{\text{total words}} - 15.59$$

Gunning Fog Index. Asses sentence complexity by analyzing polysyllabic words.

$$GFI = 0.4 \times \left(\frac{\text{total words}}{\text{total sentences}} + 100 \times \frac{\text{complex words}}{\text{total words}} \right)$$

Automated Readability Index. Counts number of characters per word and number of words per sentence.

$$ARI = 4.71 \times \frac{\text{characters}}{\text{words}} + 0.5 \times \frac{\text{words}}{\text{sentences}} - 21.43$$

3.3 Token-Probabilistic Feature Extraction

This part of the feature extraction focuses on statistical detection of machine-generated text and follows the core hypotheses of machine-generated text being more predictable than human-written text, as humans tend to write in a more surprising way, while LLMs are optimized to choose the next most likely token (Holtzman et al., 2019). To capture these metrics we extract a comprehensive set of features from the token level probability distribution generated by GPT-2.

For a given input text with the tokens $x_1, x_2, \dots, x_n$, we process the sequence with GPT-2. At the i -th position we can extract the probability distribution $P(x_{i+1} \mid x_1, \dots, x_i)$ over the entire GPT-2 Vocabulary. From this distribution we can derive different features. Here, we used features inspired by the work of Genaios, the first ranking team in the shared task and got a feature vector including the following metrics, averaged across the entire sequence.

For each position we calculate (i) the log probability of the observed token, measuring how likely the token is, given the tokens already seen. For the feature vector we then use the average observed token probability which essentially describes the perplexity of the input sequence. Perplexity is defined as the exponentiated average negative

log-likelihood of a sequence X with t tokens:

$$PPL(X) = \exp \left(-\frac{1}{t} \sum_i \log p_\theta(x_i \mid x_{<i}) \right)$$

Intuitively, it can be thought of as an evaluation of the model’s uncertainty when predicting the next token. A lower perplexity, and thus a higher average observed log-probability, is a strong indicator of machine-generated text. We further calculate (ii) the probability of the most likely token, which describes the confidence of the model at each step. Finally, we take (iii) the entropy, which measures the uncertainty of the distribution. A low entropy would indicate that the model is highly confident in a few tokens.

4 Experiments

4.1 Dataset

We used the SemEval-2024 Task 8 dataset, which is an extension of the M4 dataset (Wang et al., 2023). The training set includes five domains: Wikipedia, WikiHow, Reddit, arXiv and PeerRead. Machine-generated texts were produced by various LLMs, including davinci-003, ChatGPT, Cohere and Dolly-v2. For the test set Outfox was used as a surprising domain and GPT-4 and BLOOMz are introduced as surprising generators. The labels are roughly balanced with 56,400 machine-generated and 63,351 human-written texts. The multilingual training set additionally contains texts in Chinese, Urdu, Bulgarian and Indonesian, and for the multilingual test set, Arabic, German and Italian are introduced as unseen languages. Table 1 depicts the dataset statistics. Machine-generated texts were generated by prompting to write an article given a title (Wikipedia) or an abstract given a paper title (arXiv) and similar prompting for other domains. As in the official shared task setting, we adhered to the rules of not using additional training data, and by treating the test set labels as unavailable during the development of our model in order to assure a fair comparison to other teams results. Following the official task description, we employed accuracy as the evaluation metric to determine the quality of our model.

References

Georgios P Georgiou. 2024. Differentiating between human-written and ai-generated texts using linguistic features automatically extracted from an online computational tool. *arXiv preprint arXiv:2407.03646*.

Track	#Train	#Dev	#Test
Monolingual	119,757	5000	34,272
Multilingual	172,417	4000	42,378

Table 1: Dataset statistics for Subtask A

Yuxia Wang, Jonibek Mansurov, Petar Ivanov, Jinyan Su, Artem Shelmanov, Akim Tsvigun, Chenxi Whitehouse, Osama Mohammed Afzal, Tarek Mahmoud, Toru Sasaki, and 1 others. 2023. M4: Multi-generator, multi-domain, and multi-lingual black-box machine-generated text detection. *arXiv preprint arXiv:2305.14902*.

Zikang Guo, Kaijie Jiao, Xingyu Yao, Yuning Wan, Haoran Li, Benfeng Xu, Licheng Zhang, Quan Wang, Yongdong Zhang, and Zhendong Mao. 2024. Ustcbupt at semeval-2024 task 8: Enhancing machine-generated text detection via domain adversarial neural networks and llm embeddings. In *Proceedings of the 18th International Workshop on Semantic Evaluation (SemEval-2024)*, pages 1511–1522.

Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2019. The curious case of neural text degeneration. *arXiv preprint arXiv:1904.09751*.

Lama Ahmad Jan Hendrik Kirchner. 2023. New ai classifier for indicating ai-written text. <https://openai.com/index/new-ai-classifier-for-indicating-ai-written-text/>. Accessed on 19.08.2025.

Sarah Kreps, R Miles McCain, and Miles Brundage. 2022. All the news that’s fit to fabricate: Ai-generated text as a tool of media misinformation. *Journal of experimental political science*, 9(1):104–117.

Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D Manning, and Chelsea Finn. 2023. Detectgpt: Zero-shot machine-generated text detection using probability curvature. In *International conference on machine learning*, pages 24950–24962. PMLR.

Kseniia Petukhova, Roman Kazakov, and Ekaterina Kochmar. 2024. Petkaz at semeval-2024 task 8: Can linguistics capture the specifics of llm-generated text? *arXiv preprint arXiv:2404.05483*.

Areg Mikael Sarvazyan, José-Ángel González, and Marc Franco-Salvador. 2024. Genaios at semeval-2024 task 8: Detecting machine-generated text by mixing language model probabilistic features. In *Proceedings of the 18th International Workshop on Semantic Evaluation (SemEval-2024)*, pages 101–107.

Jan Philip Wahle, Terry Ruas, Frederic Kirstein, and Bela Gipp. 2022. How large language models are transforming machine-paraphrased plagiarism. *arXiv preprint arXiv:2210.03568*.

Yuxia Wang, Jonibek Mansurov, Petar Ivanov, Jinyan Su, Artem Shelmanov, Akim Tsvigun, Osama Mohammed Afzal, Tarek Mahmoud, Giovanni Puccetti, Thomas Arnold, and 1 others. 2024. Semeval-2024 task 8: Multidomain, multimodel and multilingual machine-generated text detection. *arXiv preprint arXiv:2404.14183*.